

# **IMAGENET ON OSG**

## **Executive summary**

We explore the feasibility and effectiveness of running ML workloads on the OSPool, with data handling mediated by OSDF/Pelican.

In this document, we limit our findings to the task of running GPU-accelerated ML training using the 150GB Imagenet input, using the recommended OSG provided method

(<https://github.com/OSGConnect/tutorial-pytorch>). The ML training job uses modest resources, using a single GPU at a time and requiring about 4h of GPU time to complete.

The input file was hosted on the SDSC OSDF/Pelican origin and was retrieved by HTCondor using the *osdf* protocol, which makes use of the OSDF redirector and the OSDF caches.

We observe that:

- It took one week to run 100 ML training jobs to completion, using about 685 GPU hours
- The success rate was about 80%, i.e. 20% of the jobs did not return a valid result
- Time spent on successful computing was about 47% of the total wallclock provided by HTCondor, with an additional 7% spent on file transfers. The file transfer overhead was thus a modest 14%.
- The significant waste was mainly due to failed (and usually re-tried) file transfers at 44%, with minimal waste due to preemption. This resulted in long tails for job completion times.
- The use of OSDF is required, as the OSPool enforces a 5GB per-job input file size limit for the native HTCondor file transfer protocol.

### The total statistics for the 100 test ML training jobs

| Description  | Total time  |
|--|-------------|
| Number of jobs that got held (with no automated recovery)  | 18          |
| Number of jobs that failed due to wrong output   | 3           |
| Total time spent computing for successful jobs   | 318.6 hours |
| Total time spent transferring files for successful jobs (last transfer, if there were retries)       | 45.8 hours  |
| Total time wasted due to transfers that were retried (for successful jobs)                           | 302.2 hours |
| Total time wasted in compute <b>due to preemption</b> (i.e. disconnects)                             | 16 hours    |
| Total time wasted in file transfers for <b>preempted jobs</b> (last transfer, if there were retries) | 0.8 hours   |
| Total time wasted in compute <b>due to wrong outputs</b>   | 0.6 hours   |
| Total time wasted in file transfer <b>for wrong outputs</b> (last transfer, if there were retries)   | 2.2 hours   |
| Mean time successful jobs took from first being matched to successful completion                     | 8.6 hours   |
| Median time successful jobs took from first being matched to successful completion                   | 4.9 hours   |

## Evaluation details

### Dataset and script used to benchmark

**Imagenet:** (150 GB) - The ImageNet dataset is used as a benchmark because its 1.2 million images across 1,000 categories test a model's ability to generalize widely. Its challenge lies in the high variability and large number of categories, making it a robust indicator of a model's capability to handle complex, real-world data.

Note: We converted the original file from zip to tar, since the OSG-recommended container image does not contain the *unzip* executable.

The ILSVRC2012\_img\_train.tar is used. **Structure of ILSVRC2012\_img\_train.tar:**

- **File Type:** tar archive containing the training images.
- **Content:** The tar file includes a directory structure where images are organized into folders corresponding to different classes.
- **Structure:**
  - **Root Directory:** The tar file typically contains a single root directory named ILSVRC2012\_img\_train.
  - **Class Subdirectories:** Inside the root directory, there are subdirectories named after each class label (e.g., n01440764, n01443537, etc.). Each subdirectory contains images belonging to that specific class.
- **Size:** The size of ILSVRC2012\_img\_train.tar is approximately **137 GB**. This large size reflects the extensive number of images included in the training set.

### **Python script explanation:**

- This Python script is designed for training a ResNet50 model on an ImageNet-like dataset using PyTorch. It begins by importing essential libraries, including `argparse` for handling command-line arguments, `torch` for deep learning operations, `torch.nn` for neural network components, and `torch.optim` for optimization algorithms. Additionally, it utilizes `torchvision` for image transformations and pre-trained models, `os` for file handling, and `PIL` (Python Imaging Library) for image processing.
- The script defines a custom dataset class, `CustomImageNetDataset`, which extends PyTorch's `Dataset` class. This class is responsible for loading and preprocessing images from a directory structure similar to ImageNet. The constructor initializes the dataset by traversing directories to collect image paths and corresponding labels. Labels are converted into integer values for model training. The `__getitem__` method loads an image, applies any specified transformations, and returns the image along with its label.
- A set of image transformations is defined to preprocess images before feeding them into the model. These transformations include resizing, horizontal flipping, tensor conversion,

and normalization, which are crucial for preparing the data in a format suitable for the ResNet50 model.

- The `ResNet50` class is defined, which wraps around PyTorch's pretrained ResNet50 model. It modifies the final fully connected layer to match the number of output classes (1000 for ImageNet). This adjustment ensures that the model can make predictions for the specific number of classes in the dataset.
- Training the model involves a `train` function that iterates over batches of data, computes the loss, performs backpropagation, and updates the model parameters using stochastic gradient descent (SGD). During training, it logs the loss and progress.
- The `main` function sets up argument parsing to handle command-line inputs, such as batch size, number of epochs, learning rate, and momentum. It initializes the dataset and data loader, sets the device to GPU if available, and creates an instance of the `ResNet50` model. The optimizer is configured with SGD and the model is trained over the specified number of epochs.
- Finally, if the `--save-model` argument is provided, the trained model's state dictionary is saved to a file named `imagenet_resnet50.pt`. This saved model can later be loaded for inference or further training.

### **Main benchmarking run**

After some initial exploration phase (described later in the document), we settled on the following job setup:

1. **Execution environment:** We used the OSG-suggested setting `+SingularityImage = "osdf:///ospool/uc-shared/public/OSG-Staff/pytorch-2.1.1.sif"`
2. **Data Fetching:** The dataset is accessed from OSDF, specifically from the path `osdf:///nrp/osdf/ILSVRC2012/ILSVRC2012_img_train.tar`.
3. **Wrapper Script Execution:** The wrapper script is responsible for setting up the environment, managing the dataset, and executing the training process. Initially, it creates a directory named `data` and moves the `ILSVRC2012_img_train.tar` file into this directory. The script then extracts the tar file, which contains the ImageNet training data.  
After extracting the primary tar file, the script proceeds to handle individual class tar files contained within. It creates directories for each class and extracts the images into their respective directories. This process involves suppressing the output of the extraction commands to avoid cluttering the logs. Finally, the script removes the tar files after extraction to save space and ensure that only the necessary files remain.
4. **Model Training:** With the data prepared, the script invokes the Python training script `main3.py`, which handles the training of the ResNet50 model. The training script is configured to save the trained model after completing the specified number of epochs. The model is trained using the PyTorch library, which is compatible with the Singularity image specified for the job.

5. **Cleanup:** After training is complete, the wrapper script performs cleanup by removing the data directory and all its contents. This ensures that no residual files are left on the compute nodes, maintaining a clean environment and freeing up storage resources

We submitted 100 jobs on the OSPool AP, each training for 5 epochs, and waited for all of them to complete.

The detailed results are listed below:

**For 5 epochs (100 jobs - submitted twice (queue 50)):**

| JOB ID         | GLIDEIN_Resource Name | TimeExecute (s)  | TimeSlot Busy (s) | Time for transferring input files (total)   | Time for last transfer (if stalled) |
|----------------|-----------------------|--|-------------------|---|-------------------------------------|
| 794569.000.000 | CHTC-Spark-CE1        | 11182  | 13283             | 35 min                                      | -                                   |
| 794569.001.000 | CHTC-Spark-CE1        | 11074  | 13177             | 35 min 2 sec                                | -                                   |
| 794569.002.000 | PDX-Coeus-CE1         | 15117  | 17833             | 45 min 10 sec                               | -                                   |
| 794569.003.000 | PDX-Coeus-CE1         | 15136  | 17872             | 45 min 29 sec                               | -                                   |
| 794569.004.000 | PDX-Coeus-CE1         | Execute: 978<br>SlotBusy: 3617<br>Input file transfer successful but no output as the network was unreachable  |                   | 13 hours 6 min 9 sec<br>(stalled 4 times)   | 43 min 55 sec                       |
| 794569.005.000 | GP-ARGO-ksu-backfill  | Job was held.<br>The job exceeded allowed execute duration of 20:00:00<br>Code 47 Subcode 0                    |                   | 13 min 19 sec                               | -                                   |
| 794569.006.000 | PDX-Coeus-CE1         | Execute: 977<br>SlotBusy: 3641<br>Input file transfer successful but no output as the network was unreachable. |                   | 11 hours 43 min 52 sec<br>(stalled 3 times) | 44 min 19 sec                       |
| 794569.007.000 | CHTC-Spark-CE1        | 11076  | 14357             | 54 min 39 sec                               | -                                   |
| 794569.008.000 | CHTC-Spark-CE1        | 11232  | 13829             | 43 min 15 sec                               | -                                   |

|                |                          |   |       |   |               |
|----------------|--------------------------|---|-------|---|---------------|
| 794569.009.000 | PDX-Coeus-CE1            | 14857   | 17558 | 44 min 55 sec   | -             |
| 794569.010.000 | PDX-Coeus-CE1            | 14892   | 17564 | 44 min 27 sec   | -             |
| 794569.011.000 | CHTC-Spark-CE1           | 10993   | 13638 | 14 hours 48 min 8 sec<br>(stalled 13 times)   | 44 min 4 sec  |
| 794569.012.000 | PDX-Coeus-CE1            | 14856   | 17850 | 8 hours 8 min 5 sec<br>(stalled 7 times)  | 49 min 48 sec |
| 794569.013.000 | CHTC-Spark-CE1           | 11506   | 13782 | 14 hours 5 min 58 sec<br>(stalled 13 times)   | 37 min 55 sec |
| 794569.014.000 | CHTC-Spark-CE1           | Output file was transferred but the model was not trained. Errors like not valid tar file and empty dataset (but the same tar file was used for all jobs) |       | 14 hours 48 min 22 sec<br>(first file transfer- 5 hours 24 min 30 sec - stalled 5 times but job was held. Tried again - 8 hours 42 min 33 sec - stalled 7 times - successful) | 41 min 54 sec |
| 794569.015.000 | CHTC-Spark-CE1           | 11323   | 11751 | 18 hours 34 min 45 sec<br>(stalled 14 times)  | 7 min 7 sec   |
| 794569.016.000 | PDX-Coeus-CE1            | 14837   | 17726 | 48 min 3 sec  | -             |
| 794569.017.000 | PDX-Coeus-CE1            | 14844   | 17829 | 49 min 39 sec   | -             |
| 794569.018.000 | GP-ARGO-ksu-b<br>ackfill | Job was held. The job exceeded allowed execute duration of 20:00:00<br>Code 47 Subcode 0  |       | 25 min 34 sec   | -             |
| 794569.019.000 | CHTC-Spark-CE1           | 11268   | 14662 | 14 hours 55 min 53 sec<br>(stalled 6 times)   | 56 min 33 sec |
| 794569.020.000 | PDX-Coeus-CE1            | 14895   | 17842 | 49 min 1 sec  | -             |
| 794569.021.000 | PDX-Coeus-CE1            | 14909   | 17601 | 2 hours 44 min 19 sec<br>(stalled 1 time)   | 44 min 46 sec |

|                |                      |       |       |  |               |
|----------------|----------------------|-------|-------|--|---------------|
| 794569.022.000 | PDX-Coeus-CE1        | 14945 | 17636 | 44 min 45 sec  | -             |
| 794569.023.000 | CHTC-Spark-CE1       | 11254 | 11722 | 7 min 46 sec   | -             |
| 794569.024.000 | GP-ARGO-ksu-backfill | 13170 | 14348 | 15 hours 19 min 29 sec<br>(stalled 13 times)   | 19 min 35 sec |
| 794569.025.000 | GP-ARGO-ksu-backfill | 16325 | 17104 | 4 hours 18 min 55 sec<br>(stalled 3 times)   | 12 min 55 sec |
| 794569.026.000 | CHTC-Spark-CE1       | 10887 | 13128 | 37 min 19 sec  | -             |
| 794569.027.000 | CHTC-Spark-CE1       | 10896 | 13157 | 37 min 38 sec  | -             |
| 794569.028.000 | PDX-Coeus-CE1        | 14893 | 17921 | 50 min 23 sec  | -             |
| 794569.029.000 | PDX-Coeus-CE1        | 14901 | 17931 | 50 min 24 sec  | -             |
| 794569.030.000 | CHTC-Spark-CE1       | 11292 | 11847 | 12 hours 24 min 9 sec<br>(first file transfer - 8 hours 31 min 42 sec -stalled 8 times but the job was held. Tried again - 2 hours 12 min 51 sec - stalled 3 times-successful) | 9 min 13 sec  |
| 794569.031.000 | CHTC-Spark-CE1       | 10891 | 12045 | 1 hour 21 min 45 sec<br>(stalled 1 time)   | 19 min 13 min |
| 794569.032.000 | CHTC-Spark-CE1       | 11475 | 12347 | 1 hour 17 min 9 sec<br>(stalled 1 time)  | 14 min 30 sec |
| 794569.033.000 | CHTC-Spark-CE1       | 11242 | 11735 | 8 min 11 sec   | -             |
| 794569.034.000 | CHTC-Spark-CE1       | 11523 | 11993 | 17 hours 32 min 10 sec<br>(first file transfer -14 hours 16 min 56 sec - stalled 12 times but the job was held. Tried again - 2 hours 12 min 20 sec                            | 7 min 49 sec  |

|                    |  |       |       |   |               |
|--------------------|--|-------|-------|---|---------------|
|                    |  |       |       | stalled 2 times -<br>successful)  |               |
| 794569.035.<br>000 | PDX-Coeus-CE1                              | 14854 | 17859 | 50 min  | -             |
| 794569.036.<br>000 | PDX-Coeus-CE1                              | 14871 | 17760 | 48 min 3 sec  | -             |
| 794569.037.<br>000 | CHTC-Spark-CE1<br>GP-ARGO-ksu-b<br>ackfill | 17608 | 18437 | 19 hours 56 min 42<br>sec<br>(first file transfer - 11<br>hours 43 min 4 sec<br>- stalled 10 times<br>but the job was<br>held. Tried again - 6<br>hours 34 min 7 sec<br>stalled 6 times but<br>the job was held<br>again. Tried again-<br>13 min 47 sec-<br>successful) | 13 min 47 sec |
| 794569.038.<br>000 | CHTC-Spark-CE1                             | 11430 | 11838 | 7 hours 33 min 33<br>sec<br>(stalled 5 times)   | 6 min 47 sec  |
| 794569.039.<br>000 | GP-ARGO-ksu-b<br>ackfill                   | 18937 | 19919 | 16 min 14 sec   | -             |
| 794569.040.<br>000 | CHTC-Spark-CE1                             | 11426 | 11886 | 4 hours 45 min 50<br>sec<br>(stalled 2 times)   | 7 min 39 sec  |
| 794569.041.<br>000 | CHTC-Spark-CE1                             | 11634 | 13907 | 1 hour 40 min 50<br>sec<br>(stalled 1 time)   | 37 min 52 sec |
| 794569.042.<br>000 | GP-ARGO-ksu-b<br>ackfill                   | 31422 | 32597 | 19 min 32 sec   | -             |
| 794569.043.<br>000 | CHTC-Spark-CE1                             | 10653 | 13698 | 50 min 42 sec   | -             |
| 794569.044.<br>000 | CHTC-Spark-CE1                             | 11089 | 13751 | 44 min 20 sec   | -             |
| 794569.045.<br>000 | CHTC-Spark-CE1                             | 11051 | 12700 | 3 hours 24 min 53<br>sec<br>(job was held once<br>and stalled 1 time)   | 27 min 28 sec |
| 794569.046.<br>000 | CHTC-Spark-CE1                             | 11119 | 11721 | 3 hours 24 min 16<br>sec<br>(job was held once)   | 10 min 1 sec  |



|                |  |   |       |  |               |
|----------------|--|---|-------|--|---------------|
| 794569.047.000 | PDX-Coeus-CE1  | 14851   | 17874 | 50 min 5 sec   | -             |
| 794569.048.000 | PDX-Coeus-CE1  | 14826   | 17845 | 50 min 14 sec  | -             |
| 794569.049.000 | CHTC-Spark-CE1   | Job 794569.049.000 was held due to a 404 error when transferring input files via the osdf protocol. |       |  |               |
| 794570.000.000 | CHTC-Spark-CE1   | 11819   | 12302 | 3 hours 22 min 34 sec<br>(stalled 1 time)  | 8 min         |
| 794570.001.000 | PDX-Coeus-CE1  | 14818   | 18265 | 57 min 21 sec  | -             |
| 794570.002.000 | PDX-Coeus-CE1  | 14853   | 18296 | 57 min 16 sec  | -             |
| 794570.003.000 | PDX-Coeus-CE1  | 14952   | 17592 | 36 hours 34 min 10 sec<br>(stalled 28 times)   | 43 min 55 sec |
| 794570.004.000 | CHTC-Spark-CE1<br>PDX-Coeus-CE1                          | 14843   | 17833 | 4 hours 37 min 37 sec (<br>- 9 min 31 sec at CHTC-Spark-CE1 but the job was evicted.<br>- 49 min 44 sec at PDX-Coeus-CE1)  | 49 min 44 sec |
| 794570.005.000 | CHTC-Spark-CE1<br>PDX-Coeus-CE1                          | 14878   | 17834 | 4 hours 11 min 15 sec (<br>- 8 min at CHTC-Spark-CE1 but the job was evicted.<br>- 49 min 10 sec at PDX-Coeus-CE1)   | 49 min 10 sec |
| 794570.006.000 | CHTC-Spark-CE1<br>CHTC-Spark-CE1<br>GP-ARGO-ksu-backfill | 24173   | 25017 | 6 hours 19 min 2 sec (<br>- 7 min 27 sec at CHTC-Spark-CE1 but the job was evicted. - 8 min 42 sec at CHTC-Spark-CE1 but the job got evicted again.<br>- 5 hours 16 min 6 sec at GP-ARGO-ksu-backfill where it stalled | 14 min 1 sec  |

|                    |   |                |                |  |               |
|--------------------|---|----------------|----------------|--|---------------|
|                    |   |                |                | once)  |               |
| 794570.007.<br>000 | GP-ARGO-ksu-b<br>ackfill  | 19034          | 19965          | 12 hours 5 min 37<br>sec<br>(stalled 1 time)                                     | 15 min 26 sec |
| 794570.008.<br>000 | PDX-Coeus-CE1   | 14587          | 16796          | 36 min 44 sec  | -             |
| 794570.009.<br>000 | PDX-Coeus-CE1   | 14603          | 16810          | 36 min 42 sec  | -             |
| 794570.010.<br>000 | GP-ARGO-ksu-b<br>ackfill  | 19746<br>49304 | 20691<br>50091 | 6 hours 42 sec (15<br>min 39 sec first file<br>transfer but job got<br>evicted ) | 13 min 4 sec  |
| 794570.011.<br>000 | PDX-Coeus-CE1   | 14673          | 17078          | 39 min 59 sec  | -             |
| 794570.012.<br>000 | PDX-Coeus-CE1   | 14620          | 17072          | 40 min 45 sec  | -             |
| 794570.013.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.014.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.015.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.016.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.017.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.018.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.019.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.020.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.021.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.022.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |
| 794570.023.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |                |                |  |               |


|                    |   |       |       |   |               |
|--------------------|---|-------|-------|---|---------------|
| 794570.024.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |       |       |   |               |
| 794570.025.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |       |       |   |               |
| 794570.026.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |       |       |   |               |
| 794570.027.<br>000 | Job was held due to a failure in transferring input files, with a "404 page not found" error. |       |       |   |               |
| 794570.028.<br>000 | PDX-Coeus-CE1   | 14967 | 17620 | 11 hours 40 min 10<br>sec<br>(stalled 8 times)  | 44 min 8 sec  |
| 794570.029.<br>000 | PDX-Coeus-CE1   | 14609 | 16769 | 35 min 54 sec                                   | -             |
| 794570.030.<br>000 | PDX-Coeus-CE1   | 14635 | 16828 | 36 min 28 sec                                   | -             |
| 794570.031.<br>000 | PDX-Coeus-CE1   | 14815 | 17846 | 50 min 25 sec                                   | -             |
| 794570.032.<br>000 | PDX-Coeus-CE1   | 14848 | 17867 | 50 min 14 sec                                   | -             |
| 794570.033.<br>000 | PDX-Coeus-CE1   | 14835 | 17527 | 44 min 47 sec                                   | -             |
| 794570.034.<br>000 | PDX-Coeus-CE1   | 14862 | 17560 | 44 min 52 sec                                   | -             |
| 794570.035.<br>000 | PDX-Coeus-CE1   | 14726 | 16902 | 17 hours 16 min 10<br>sec<br>(stalled 13 times) | 36 min 11 sec |
| 794570.036.<br>000 | PDX-Coeus-CE1   | 14659 | 16785 | 17 hours 25 min 20<br>sec<br>(stalled 13 times) | 35 min 20 sec |
| 794570.037.<br>000 | GP-ARGO-ksu-b<br>ackfill  | 58683 | 59669 | 16 min 23 sec                                   | -             |
| 794570.038.<br>000 | PDX-Coeus-CE1   | 14621 | 17040 | 40 min 14 sec                                   | -             |
| 794570.039.<br>000 | PDX-Coeus-CE1   | 14586 | 16991 | 39 min 59 sec                                   | -             |
| 794570.040.<br>000 | GP-ARGO-ksu-b<br>ackfill  | 14598 | 15362 | 12 min 37 sec                                   | -             |

|                    |                          |       |       |  |               |
|--------------------|--------------------------|-------|-------|--|---------------|
| 794570.041.<br>000 | PDX-Coeus-CE1            | 14821 | 17488 | 15 hours 4 min 20<br>sec<br>(stalled 13 times) | 44 min 20 sec |
| 794570.042.<br>000 | PDX-Coeus-CE1            | 14841 | 17508 | 15 hours 4 min 20<br>sec<br>(stalled 13 times) | 44 min 21 sec |
| 794570.043.<br>000 | PDX-Coeus-CE1            | 14947 | 17570 | 43 min 38 sec                                  | -             |
| 794570.044.<br>000 | PDX-Coeus-CE1            | 14872 | 17407 | 41 min 56 sec                                  | -             |
| 794570.045.<br>000 | GP-ARGO-ksu-b<br>ackfill | 14324 | 15088 | 12 min 40 sec                                  | -             |
| 794570.046.<br>000 | GP-ARGO-ksu-b<br>ackfill | 14154 | 14965 | 13 min 27 sec                                  | -             |
| 794570.047.<br>000 | PDX-Coeus-CE1            | 14808 | 17792 | 49 min 38 sec                                  | -             |
| 794570.048.<br>000 | PDX-Coeus-CE1            | 14840 | 17827 | 49 min 41 sec                                  | -             |
| 794570.049.<br>000 | GP-ARGO-ksu-b<br>ackfill | 12467 | 13300 | 13 min 51 sec                                  | -             |

Time successful jobs took from first being matched to successful completion (starting file transfer to job termination) :



The line chart depicts the time taken for successful job completion from the start of file transfer to job termination across multiple job IDs. The completion times exhibit significant variability, with most jobs completing between 20,000 and 60,000 seconds. However, there are noticeable spikes, with some jobs taking over 100,000 seconds, indicating occasional outliers that experience much longer processing times. These outliers suggest variability in resource allocation or job complexity.




The bar chart provides a comparison of mean and median job compute times across three sites: CHTC-SpArk-CE1, GP-ARGO-Xsu-backfill, and PDX-Coeus-CE1.

- The data reveals that GP-ARGO-Xsu-backfill has significantly higher compute times, with a mean of 23,400 seconds and a median of 17,608 seconds, indicating a larger variability in job durations at this site.
- In contrast, CHTC-SpArk-CE1 exhibits the lowest compute times, with a mean of 11,222 seconds and a median of 11,237 seconds, suggesting a consistent job duration with minimal variability.
- PDX-Coeus-CE1 falls in the middle, with mean and median compute times close to each other, at 14,824 seconds and 14,846 seconds, respectively, indicating a consistent performance despite handling the largest number of jobs (42).

GP-ARGO-Xsu-backfill experiences higher compute times, possibly due to factors like resource availability or job complexity.

CHTC-SpArk-CE1 and PDX-Coeus-CE1 have more consistent and lower compute times, with PDX-Coeus-CE1 managing a larger workload effectively.



The bar chart illustrates the time wasted due to stalling for each job ID, showing significant variation across jobs. Some jobs experience minimal stalling, while others waste substantial time, with one job wasting over 129,000 seconds. The high variability in stalling times suggests inefficiencies in resource allocation or scheduling, leading to delays in job processing. This highlights the need for optimization to reduce wasted time and improve overall job efficiency.

**Unsuccessful jobs:**

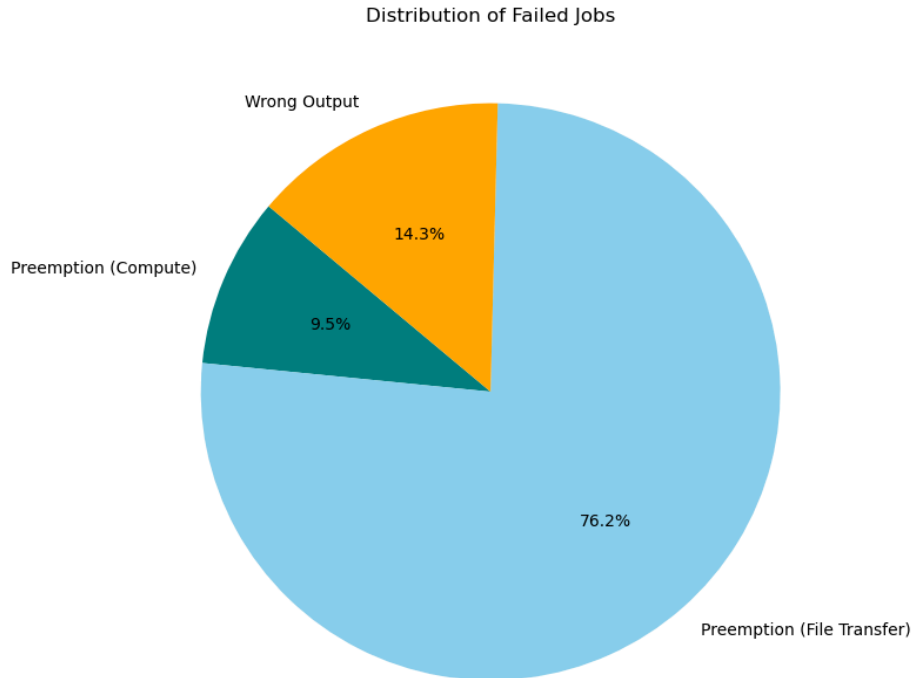
Failed jobs (wrong output):

| JOB ID         | TimeExecute (s) | Time for last transfer<br>(if stalled) |
|----------------|-----------------|--|
| 794569.004.000 | 978             | 43 min 55 sec                          |
| 794569.006.000 | 977             | 44 min 19 sec                          |
| 794569.014.000 | 33              | 41 min 54 sec                          |

Preempted failed job (job was held):

| JOB ID         | Job Held during: | TimeExecute (s) | Time for last<br>transfer<br>(if stalled) |
|----------------|------------------|-----------------|---|
| 794569.005.000 | Compute          | 72024           | 13 min 19 sec                             |
| 794569.018.000 | Compute          | 72065           | 25 min 34 sec                             |
| 794569.049.000 | File transfer    | -               | 1 sec                                     |
| 794570.013.000 | File transfer    | -               | 1 sec                                     |
| 794570.014.000 | File transfer    | -               | 1 sec                                     |
| 794570.015.000 | File transfer    | -               | 1 sec                                     |
| 794570.016.000 | File transfer    | -               | 1 sec                                     |
| 794570.017.000 | File transfer    | -               | 1 sec                                     |
| 794570.018.000 | File transfer    | -               | 1 sec                                     |
| 794570.019.000 | File transfer    | -               | 1 sec                                     |
| 794570.020.000 | File transfer    | -               | 1 sec                                     |
| 794570.021.000 | File transfer    | -               | 1 sec                                     |
| 794570.022.000 | File transfer    | -               | 1 sec                                     |
| 794570.023.000 | File transfer    | -               | 1 sec                                     |
| 794570.024.000 | File transfer    | -               | 1 sec                                     |
| 794570.025.000 | File transfer    | -               | 1 sec                                     |
| 794570.026.000 | File transfer    | -               | 1 sec                                     |
| 794570.027.000 | File transfer    | -               | 1 sec                                     |





The chart provides a detailed analysis of the reasons behind job failures within a particular system or process. The distribution is divided into three primary categories:

- **Preemption (File Transfer):** This category represents the largest share of failures, accounting for 76.2% of the unsuccessful jobs. Preemption here likely refers to interruptions or terminations of file transfer processes before they could be completed. This could be due to resource reallocation, system priorities, or other external factors that caused the file transfers to fail.
- **Wrong Output:** This is the second most common cause of failure, making up 14.3% of the unsuccessful jobs. Jobs categorized under "Wrong Output" are those that completed execution but produced incorrect or unexpected results, which would require reprocessing or debugging to correct.
- **Preemption (Compute):** This category, which accounts for 9.5% of the failures, involves jobs that were preempted during the computation phase. Similar to file transfer preemption, these interruptions could have been caused by resource shortages, higher-priority processes, or system errors that halted the computation mid-process.

The total number of failed jobs is 21, indicating that these 21 jobs did not complete successfully for the reasons mentioned above. Despite these failures, the overall success rate of the jobs stands at 79%, suggesting that the majority of jobs were completed without issues.

**Total time statistic for 100 jobs (5 epochs):**

Number of unsuccessful jobs: 21

**Success rate: 79%**

| Description  | Total time                     |
|--|--------------------------------|
| Total time spent computing for successful jobs   | 13 days 6 hours 33 min 53 sec  |
| Total time spent transferring files for successful jobs (last transfer, if there were retries)                               | 1 day 21 hours 50 min 43 sec   |
| Total time wasted due to transfers that failed (for successful jobs)   | 12 days 14 hours 12 min 50 sec |
| Number of jobs that failed (with no automated recovery)  | 18                             |
| Number of jobs that failed due to wrong output   | 3                              |
| Total time wasted in compute <b>due to preemption</b> (i.e. disconnects)   | 16 hours 1 min 29 sec          |
| Total time wasted in file transfers for <b>preempted jobs</b> (last transfer, if there were retries)                         | 47 min 35 sec                  |
| Total time wasted in compute <b>due to wrong outputs</b>   | 33 min 8 sec                   |
| Total time wasted in file transfer <b>for wrong outputs</b> (last transfer, if there were retries)                           | 2 hours 11 min 8 sec           |
| Total time wasted in compute <b>due ALL failed jobs</b>  | 16 hours 34 min 37 sec         |
| Total time wasted in file transfer <b>for ALL failed jobs</b> (last transfer, if there were retries)                         | 2 hours 58 min 43 sec          |
| Mean time successful jobs took from first being matched to successful completion (starting file transfer to job termination) | 8 hours 36 min 46 sec          |
| Median time successful jobs took from first being matched to successful completion   | 4 hours 57 min 25 sec          |

Across the 100 jobs, the majority were run on three sites: CHTC-Spark-CE1, PDX-Coeus-CE1, and GP-ARGO-ksu-backfill. The jobs were divided into two queues of 50 jobs each.

### **Job Execution Performance Analysis**

The success rate of the jobs was 79%, indicating that a significant majority of the jobs were completed successfully. For these successful jobs, the total time spent on computing was 13 days, 6 hours, 33 minutes, and 53 seconds. Additionally, file transfers for these jobs took 1 day, 21 hours, 50 minutes, and 43 seconds (the last transfer if there were retries). However, there

was considerable time wasted due to failed file transfers, amounting to 12 days, 14 hours, 12 minutes, and 50 seconds, despite these jobs ultimately being successful.

In terms of job failures, 18 jobs did not recover automatically, and 3 of these failures were due to incorrect output. The time wasted in computing due to preemption—such as job disconnects—was 16 hours, 1 minute, and 29 seconds. File transfers for preempted jobs added an additional 47 minutes and 35 seconds of wasted time. Moreover, the total time wasted in computing due to wrong outputs was 33 minutes and 8 seconds, with file transfers for these wrong outputs contributing another 2 hours, 11 minutes, and 8 seconds. Overall, the total time wasted in computing for all failed jobs was 16 hours, 34 minutes, and 37 seconds, while the time wasted in file transfers amounted to 2 hours, 58 minutes, and 43 seconds.

Regarding performance metrics, the mean time from when a job was first matched to its successful completion, including file transfer and job termination, was 8 hours, 36 minutes, and 46 seconds. The median time for successful jobs was 4 hours, 57 minutes, and 25 seconds.

Overall, while most jobs succeed, the system loses considerable time due to failed transfers, preemptions, and incorrect outputs, indicating areas for potential improvement.

## **Initial tries and learning path**

Training models on the ImageNet dataset presents significant challenges due to the dataset's large size and complexity. To efficiently utilize computational resources on the Open Science Grid (OSG), effective strategies for managing and processing this large dataset are crucial. The goal is to benchmark various methods for handling the dataset, evaluating their efficiency, resource usage, and overall feasibility.

In this context, three methods have been explored for training ImageNet on OSG:

1. **Method 1: Traditional Download and Extraction with HTCondor native Handling** - This approach involves downloading the dataset directly to the AP, use the native HTCondor file transfer, and then using it for model training. This method is straightforward but requires substantial local storage and can be constrained by HTCondor's file transfer limits.
2. **Method 2: Direct Fetching directly from the Origin with HTTP** - This method simplifies the process by fetching the dataset directly from the origin via HTTPS within the job script. It integrates dataset retrieval, extraction, training, and cleanup into a single script, reducing the need for pre-job data management.
3. **Method 3: Direct Fetching from the Origin Using OSDF Protocol** - This approach leverages the OSDF protocol to efficiently fetch the dataset, streamlining data handling and reducing AP storage requirements. It is designed to efficiently manage large datasets by utilizing the capabilities of OSDF for better data transfer and access.

Each method offers a different approach to managing the dataset, and their effectiveness in training ImageNet on OSG will be evaluated based on their resource utilization, efficiency, and practicality.

## **METHOD 1:**

### **Introduction:**

Method 1 involves a straightforward approach where the dataset is downloaded locally on the AP, transferred to the computed node by HTCondor, and then used for training the model.

### **Result:**

During the implementation of Method 1, an error was encountered due to the size of the input files exceeding the maximum allowed transfer limit configured in HTCondor. Specifically, the job was held with the following error message:

```
2024-07-18 14:35:31 Job was held.
```

```
    TransferInputSizeMB (158783) is greater than MAX_TRANSFER_INPUT_MB  
(5000) at submit time  
    Code 32 Subcode 0
```

This error indicates that the total size of the input files (approximately 150 GB) surpassed the maximum transfer limit of 5000 MB (5 GB). HTCondor imposes this limit to manage resource usage and ensure efficient job scheduling, which led to the job being held.

We attempted splitting the large input file in many smaller chunks, and list them separately in the job description. Even though the dataset was divided into chunks, the total size of the input files reported was still above the 5 GB threshold. This is because HTCondor has a hard limit of 5 GB for the maximum input size per job, and it cannot handle input files exceeding this limit, regardless of how they are split.

Method 1 thus ultimately proved **unsuccessful**.

## **METHOD 2:**

### **Introduction:**

Method 2 involves accessing the dataset directly from OSDF via HTTP, bypassing the need for manual downloading of the files to the AP. The workflow for Method 2 includes direct access to the dataset, executing the model training process, and performing cleanup operations, all managed through HTCondor for job submission and resource allocation.

### **Process:**

1. **Data Fetching:** The dataset is accessed directly from the provided HTTPS URL: (transfer\_input\_files=https://sdsc-origin.nationalresearchplatform.org:8443/nrp/osdf/ILSVRC2012/ILSVRC2012\_img\_train.tar). This approach uses the HTTPS protocol for secure data transfer, avoiding the need for manual downloading or extraction of tarballs.
2. **Wrapper Script Execution:**
  - The wrapper script sets up the environment and prepares the dataset for training. It begins by creating a directory named data and then moves the ILSVRC2012\_img\_train.tar file into this directory. The script then extracts the tar file containing the ImageNet training data.
  - After extracting the primary tar file, the script handles any additional tar files found within. It creates directories for each class and extracts the images into their respective directories. To keep the logs clean, the extraction output is suppressed. Finally, the script removes the tar files after extraction to save space and ensure that only the necessary files remain.
3. **Model Training:** With the data prepared, the script executes the Python training script main2.py. This script trains the ResNet50 model using PyTorch, with the training configured to save the model after the specified number of epochs. The use of PyTorch and the Singularity image ensures compatibility and efficient execution.
4. **Cleanup:** After the training process is complete, the wrapper script cleans up by removing the data directory and all its contents. This step ensures that no residual files are left on the compute nodes, maintaining a clean environment and freeing up storage resources.

### Result:

During the implementation of Method 2, an error was encountered. The transfer failed with an error indicating a lack of progress, which led to job eviction and holding. This issue is likely due to network problems, server unresponsiveness, or transfer interruptions.

012 (791469.000.000) 2024-08-09 22:31:42 Job was held.

Transfer input files failure at the execution point using protocol https. Details: Aborted due to lack of progress

The file transfer was interrupted because it was not making adequate progress. The transfer from the specified URL to the job's execution node failed, possibly due to network issues, slow download speeds, or server-side problems. The system aborted the transfer after it determined that progress was insufficient, preventing the job from starting successfully.

During an attempt to rerun, the same error occurred. Further using wget in the wrapper script was tried.

**Proposed Solution - wget command in wrapper script** - In this method, the wrapper script is configured to use wget to fetch the ImageNet dataset directly from OSDF via HTTP. By

specifying the URL, the script automates the process of downloading the dataset during job execution.

Unfortunately, wget is not installed in the used singularity image, so we skipped this evaluation completely.

## **METHOD 3:**

### **Introduction:**

Method 3 involves fetching the dataset directly from the origin using the Open Science Data Framework (OSDF) protocol. This approach is designed to streamline data handling by leveraging OSDF's capabilities for efficient data transfer and access. The dataset is accessed through OSDF during the job execution, which reduces the need for substantial local storage and minimizes the overhead associated with data management. By integrating dataset retrieval directly within the job script, Method 3 simplifies the workflow and ensures that large datasets can be handled effectively without overwhelming local storage resources. This method is particularly advantageous in environments where managing large volumes of data efficiently is critical.

### **Process:**

Method 3 involves leveraging the Open Science Data Framework (OSDF) protocol to efficiently manage and process the ImageNet dataset during model training. This approach is designed to streamline data handling by accessing the dataset directly from OSDF, thereby reducing the need for extensive local storage and simplifying the workflow.

### **Detailed Steps:**

6. **Data Fetching:** The dataset is accessed from OSDF, specifically from the path `osdf:///nlp/osdf/ILSVRC2012/ILSVRC2012_img_train.tar`. This method utilizes OSDF's efficient data transfer capabilities, which are designed to handle large datasets and minimize the overhead of manual data management.
7. **Wrapper Script Execution:** The wrapper script is responsible for setting up the environment, managing the dataset, and executing the training process. Initially, it creates a directory named `data` and moves the `ILSVRC2012_img_train.tar` file into this directory. The script then extracts the tar file, which contains the ImageNet training data.  
After extracting the primary tar file, the script proceeds to handle individual class tar files contained within. It creates directories for each class and extracts the images into their respective directories. This process involves suppressing the output of the extraction commands to avoid cluttering the logs. Finally, the script removes the tar files after extraction to save space and ensure that only the necessary files remain.
8. **Model Training:** With the data prepared, the script invokes the Python training script `main3.py`, which handles the training of the ResNet50 model. The training script is configured to save the trained model after completing the specified number of epochs.

The model is trained using the PyTorch library, which is compatible with the Singularity image specified for the job.

9. **Cleanup:** After training is complete, the wrapper script performs cleanup by removing the data directory and all its contents. This ensures that no residual files are left on the compute nodes, maintaining a clean environment and freeing up storage resources

## **Errors:**

### **1. SciPy Module Installation Error**

During the job execution, an `OSError: [Errno 38] Function not implemented` was encountered while installing the SciPy module. This error typically occurs due to limitations or restrictions in the filesystem or environment permissions, which might prevent certain operations from being performed. Additionally, a `ModuleNotFoundError` indicated that the SciPy module was not available when the script attempted to use it. To address this, manual installation of the SciPy module using `pip install scipy` was done. This action resolved the missing module issue, allowing the job to continue its execution without encountering further errors related to SciPy.

### **2. Failed File Transfer (404 Error)**

The job initially failed to transfer a file from the URL `osdf:///ospool/imagenet/imagenet-object-localization-challenge.zip`, resulting in a 404 error. This error indicates that the file or the specified path could not be found on the storage system. Such issues can arise from incorrect file paths, non-existent files, or misconfigured storage systems. After identifying the problem, the transfer process was retried, which was successful on subsequent attempts. This allowed the necessary file to be transferred and made available for the job, resolving the initial file transfer failure.

### **3. Zip File Handling Failure**

The job encountered issues with handling a zip file due to the absence of the `unzip` command in the execution environment. This problem prevented the extraction of data from the zip file, which was crucial for the job. To overcome this issue, the problematic zip file was replaced with a tar file. The tar format was successfully processed and extracted without encountering the same issues, thus enabling the job to access and use the data as intended. This change resolved the difficulties related to zip file handling and ensured smooth data processing.

### **4. Archive Problem**

A problem arose when the job attempted to access the dataset archive `ILSVRC2012_devkit_t12.tar.gz`. This file was missing, leading to runtime errors in the Python script that relied on it for data processing. The absence of this archive prevented the job from proceeding as expected. As a corrective measure, the dataset archive, `ILSVRC2012_img_train.tar` was continued to be used, which was available for training the

model by creating a custom dataset class, CustomImageNetDataset, to handle the training data more effectively. This adjustment allowed the training process to move forward without being hindered by the corrupted or missing file.

## **5. Socket Connection Closure and Job Disconnection**

The job experienced disconnection issues due to unexpected socket closures between the submit and execution hosts. This led to the job exceeding the 40-minute reconnection window, which is a common timeout limit in job scheduling systems. Network disruptions or problems with the execution host can cause such issues. To mitigate the risk of disconnections, the allocated disk space and other resources for the job were increased. This adjustment was aimed at providing a more stable environment and reducing the likelihood of similar disconnection problems in future job executions.

## **6. Repeated Disconnections and Job Lease Expiration**

The job faced multiple disconnections, which led to lease expiration and eventual eviction. Lease expiration occurs when a job cannot maintain a connection within the allocated time frame, resulting in its eviction from the queue. This problem is often due to persistent network or system failures. To address this issue, the allocated disk space and other resources were increased for the job, which helped to stabilize the environment and minimize the likelihood of further disconnections. These adjustments improved the job's stability, allowing it to complete successfully in subsequent runs.

## **7. Persistent File Transfer Stalls and Job Evictions**

During the job execution, there were repeated stalls in the file transfer process, leading to multiple job evictions. The stalls in file transfers and subsequent job evictions were primarily due to an overloaded cache, which impeded the efficient handling of large files. When the cache becomes overloaded, it can cause significant delays in file transfer operations, leading to interruptions in job execution. Network instability may also exacerbate these issues, but the primary cause in this case was the cache overload. Once the cache-related issues were addressed, the job was able to execute without further errors. They still stall sometimes now but the file transfer is successful after some tries.

### **Result:**

The ResNet model was successfully trained on the ImageNet dataset for 20 epochs, with the entire process taking approximately 14 hours. The training was completed on the resource `GLIDEIN_ResourceName = "GP-ARGO-ksu-backfill"`. After overcoming the initial file transfer and execution issues, the job was executed without further errors and finished successfully.

Subsequently, the training process was tested by queuing the job five times for 20 epochs. These jobs were distributed across different sites, demonstrating the model's capability to run



consistently across multiple resources. The tests validated the robustness and portability of the training setup, confirming successful execution in diverse environments.

**For 20 epochs:**

| JOB ID         | GLIDEIN_Resource Name                           | TimeExecute (s) | TimeSlotBusy (s) | Time for transferring input files (total)                   |
|----------------|---|-----------------|------------------|---|
| 791459.000.000 | PDX-Coeus-CE1                                   | 56290           | 57782            | 7 hours, 42 min, and 26 sec (file transfer stalled 7 times) |
| 791459.001.000 | Stuck job, kept stalling. It was removed later. |                 |                  |   |
| 791459.002.000 | CHTC-Spark-CE1                                  | 43339           | 46489            | 53 min and 28 sec   |
| 791459.003.000 | GP-ARGO-ksu-backfill                            | 64328           | 65174            | 14 min and 2 sec  |
| 791459.004.000 | PDX-Coeus-CE1                                   | 56057           | 57795            | 5 hours, 56 min, and 1 sec (file transfer stalled thrice)   |

The job, which involved training a ResNet50 model on the ImageNet dataset, was successfully executed across several sites. The following observations and results were noted:

**1. Successful Executions:**

- Jobs running on PDX-Coeus-CE1, CHTC-Spark-CE1, and GP-ARGO-ksu-backfill were completed successfully. Each site demonstrated varying performance in terms of file transfer and overall execution time.

**2. Performance Insights:**

- Least Time for Transferring Input Files:** The job executed on GP-ARGO-ksu-backfill exhibited the shortest file transfer time, taking only 14 minutes and 2 seconds. This indicates efficient handling of file transfers at this site.
- Least Time to Execute:** The job on CHTC-Spark-CE1 had the shortest overall execution time, completing in 12 hours. This reflects efficient processing and resource utilization during model training.

**3. Overall Efficiency:**

- The job's success across different sites highlights the importance of both effective file transfer and computational efficiency. GP-ARGO-ksu-backfill was notably efficient in file transfers, while CHTC-Spark-CE1 achieved the best overall execution time.

- Additionally, it was observed that when input files are transferred for the second time on the same site, the process takes significantly less time, leading to faster results. This indicates that once a site has preloaded the necessary data, subsequent jobs benefit from reduced transfer times, further enhancing overall job efficiency.

**For 5 epochs:**

| JOB ID         | GLIDEIN_Resource Name | TimeExecute (s) | TimeSlotBusy (s) | Time for transferring input files (total)                              |
|----------------|-----------------------|-----------------|------------------|--|
| 794421.000.000 | CHTC-Spark-CE1        | 10209           | 10932            | 12 min 1 sec   |
| 794421.001.000 | GP-ARGO-ksu-backfill  | 16757           | 17539            | 12 min 58 sec  |
| 794421.002.000 | CHTC-Spark-CE1        | 11039           | 11533            | 1 hour 8 min 41 sec (stalled once);<br>last transfer- 8 min 12 sec     |
| 794421.003.000 | GP-ARGO-ksu-backfill  | 17424           | 18202            | 2 hours 52 min 39 sec (stalled once) ;<br>last transfer- 12 min 53 sec |
| 794421.004.000 | CHTC-Spark-CE1        | 10407           | 11132            | 12 min 3 sec   |

The jobs executed across CHTC-Spark-CE1 and GP-ARGO-ksu-backfill sites yielded important insights into the performance and efficiency of the model's training process.

1. **Consistency in Execution Times:** Jobs run on CHTC-Spark-CE1 exhibited more consistent and slightly shorter execution times compared to GP-ARGO-ksu-backfill. Specifically, the job on CHTC-Spark-CE1 with Job ID 794421.000.000 completed in 10,209 seconds (approximately 2 hours 50 minutes), while GP-ARGO-ksu-backfill took 16,757 seconds (approximately 4 hours 40 minutes) for Job ID 794421.001.000.
2. **Impact of File Transfer Delays:** The transfer times varied significantly across different runs. The first job on GP-ARGO-ksu-backfill had a reasonable transfer time of 12 minutes 58 seconds, but a later job (794421.003.000) on the same site faced a substantial delay, with the transfer taking 2 hours 52 minutes and 39 seconds due to a stall. The same happened with CHTC-Spark-CE1.

Method 3 effectively utilized the Open Science Data Framework (OSDF) to streamline the ImageNet dataset handling and model training process.

Subsequent tests across different sites confirmed the method's robustness:

- **Shortest File Transfer Time:** GP-ARGO-ksu-backfill at 14 minutes and 2 seconds.
- **Fastest Overall Execution Time:** CHTC-Spark-CE1 with a 12-hour total runtime.
- **Failures (Stuck and Other Issues):** 20% (takes a lot of time to complete)
- **Success Rate:** 80%

These results suggest that selecting sites with efficient file handling capabilities and optimized computational resources can significantly impact job performance and completion time.

In conclusion for 5 epochs, while both CHTC-Spark-CE1 and GP-ARGO-ksu-backfill demonstrated their capabilities, the occurrence of stalling significantly impacted the efficiency of job executions. If stalling during file transfers can be avoided or minimized, the overall performance and efficiency of both sites could be greatly improved, resulting in more reliable and faster job completions.

## **Overall Conclusion:**

Training the ResNet50 model on the ImageNet dataset across different methods on the Open Science Grid (OSG) provided valuable insights into dataset handling and job execution efficiency.

**Method 1: Traditional Download and Extraction** This method was unsuccessful due to HTCondor's 5 GB input file transfer limit. Despite efforts to split the dataset into smaller chunks, the total size of the input files consistently exceeded the limit, preventing successful execution.

**Method 2: Direct Fetching from Origin with HTTP** Method 2 also faced challenges. The initial file transfer issues led to job eviction due to lack of progress, likely caused by network or server problems. Attempts to use `wget` for downloading the dataset failed due to the unavailability of the command in the execution environment. Consequently, this method was deemed unsuccessful.

**Method 3: Direct Fetching from Origin Using OSDF Protocol** Method 3 proved successful, leveraging the Open Science Data Framework (OSDF) to efficiently manage the ImageNet dataset. After overcoming initial issues with file transfers and environment configurations, the ResNet50 model was trained for 20 epochs successfully. Testing across various sites demonstrated the robustness of this method. The 5 epoch runs showed that stalling during file transfers posed a challenge that, if mitigated, would significantly enhance overall efficiency.

Overall, Method 3 was the most effective approach for handling large datasets on OSG, highlighting the importance of efficient data transfer protocols and site-specific performance optimization. This method demonstrated the ability to handle large datasets effectively while ensuring robust model training across different environments. Addressing the stalling issues observed in file transfers would make this approach even more reliable and efficient for large-scale distributed training on the OSG.

**Future Scope:**

In summary, the key areas for improvement are optimizing file transfer processes to minimize inefficiencies, enhancing automated recovery and failure management to reduce the impact of job failures, and striving for more consistent job completion times. By focusing on these areas, the system's overall performance and reliability can be significantly improved.